

```

import tensorflow as tf
from tensorflow.keras import models, layers
import matplotlib.pyplot as plt

from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

IMAGE_SIZE = 256
BATCH_SIZE = 32
CHANNELS = 3
EPOCHS = 50

dataset = tf.keras.preprocessing.image_dataset_from_directory(
    "/content/drive/MyDrive/PlantVillage",
    shuffle = True,
    image_size = (IMAGE_SIZE,IMAGE_SIZE),
    batch_size = BATCH_SIZE
)

Found 2152 files belonging to 3 classes.

class_names = dataset.class_names
class_names

['Potato__Early_blight', 'Potato__Late_blight', 'Potato__healthy']

len(dataset)

68

for image_batch, label_batch in dataset.take(1):
    print(image_batch.shape)
    print(label_batch.numpy())

(32, 256, 256, 3)
[0 1 1 0 1 0 1 0 1 0 1 0 1 0 0 1 2 1 0 0 0 0 1 0 1 0 2 1 1]

for image_batch, label_batch in dataset.take(1):
    # without .numpy() function it give tensor

    print(image_batch[0].numpy())
    # image_batch[0].numpy() it gives 3D array

    # these values are float we have to convert them into int using astype() function
    plt.imshow(image_batch[0].numpy().astype("uint8"))
    plt.axis("off")

```

## Resources X

...

You are not subscribed. [Learn more](#)

You currently have zero compute units available.

Resources offered free of charge are not guaranteed.

Purchase more units [here](#).

At your current usage level, this runtime may last up to 4 hours.

[Manage sessions](#)

Want more memory and disk space? X

[Upgrade to Colab Pro](#)

Python 3 Google Compute Engine backend (GPU)

Showing resources from 1:43 AM to 1:55 AM

System RAM  
7.0 / 12.7 GB



GPU RAM  
6.2 / 15.0 GB



Disk  
37.4 / 112.6 GB



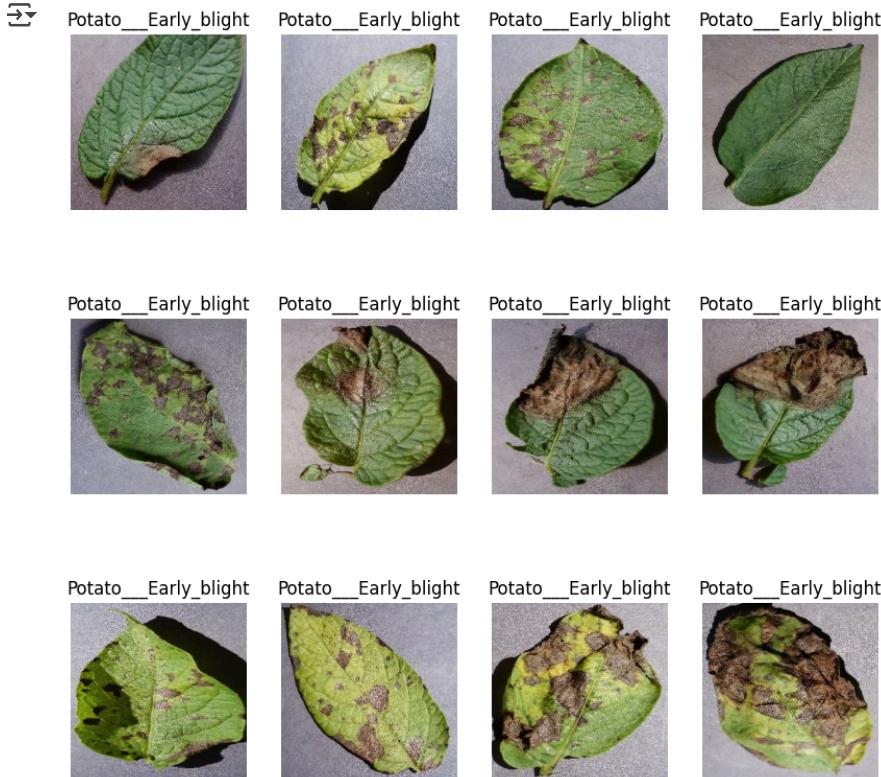
```
→ [[[189. 180. 183.]  
[190. 181. 184.]  
[193. 184. 187.]  
...  
[162. 155. 162.]  
[159. 152. 159.]  
[157. 150. 157.]]  
  
[[186. 177. 180.]  
[186. 177. 180.]  
[187. 178. 181.]  
...  
[161. 154. 161.]  
[159. 152. 159.]  
[158. 151. 158.]]  
  
[[184. 175. 178.]  
[182. 173. 176.]  
[182. 173. 176.]  
...  
[159. 152. 159.]  
[159. 152. 159.]  
[159. 152. 159.]]  
  
...  
  
[[151. 140. 138.]  
[145. 134. 132.]  
[159. 148. 146.]  
...  
[110. 99. 105.]  
[127. 116. 122.]  
[143. 132. 138.]]  
  
[[153. 142. 140.]  
[150. 139. 137.]  
[156. 145. 143.]  
...  
[109. 98. 104.]  
[127. 116. 122.]  
[138. 127. 133.]]  
  
[[155. 144. 142.]  
[153. 142. 140.]  
[144. 133. 131.]  
...  
[120. 109. 115.]  
[125. 114. 120.]  
[119. 108. 114.]]]
```



```
# plt.figure(figsize = (10,10)) WE USED IT SO THE IMAGES DONT OVERLAP BELOW  
plt.figure(figsize = (10,10))  
for image_batch, label_batch in dataset.take(1):
```

```
# we onnly took 12 images from teh dataset to print beacuse 32 will ake it diff
for i in range(12):
    # subplot help us to plot multiple images at a time
    ax = plt.subplot(3,4,i+1)

    # without .numpy() functio it give tensor
    # these values are float we have to convert them into int using astype() fu
    plt.imshow(image_batch[i].numpy().astype("uint8"))
    plt.title(class_names[label_batch[1]])
    plt.axis("off")
```



```
len(dataset)
```

→ 68

**80% ==> training**

**20% ==> 10% validation, 10% testing**

```
def get_dataset_partitions_tf(ds, train_split=0.8, val_split=0.1, test_split=0.1,
                             assert (train_split + test_split + val_split) == 1

    ds_size = len(ds)

    if shuffle:
        # seed is just for predictibilty, if you do same seed everytime you get s
        ds = ds.shuffle(shuffle_size, seed=12)

    train_size = int(train_split * ds_size)
    val_size = int(val_split * ds_size)
```

```

train_ds = ds.take(train_size)
val_ds = ds.skip(train_size).take(val_size)
test_ds = ds.skip(train_size).skip(val_size)

return train_ds, val_ds, test_ds

train_ds, val_ds, test_ds = get_dataset_partitions_tf(dataset)

len(train_ds)
→ 54

len(val_ds)
→ 6

len(test_ds)
→ 8

# cache() it will read the image from the disk and then for the next iteration it
# prefetch(), if you are using GPU and CPU, if GPU is busy training prefetch will
# buffer_size = tf.data.AUTOTUNE using this i am letting tensorflow determine how

train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size = tf.data.AUTOTUNE)
val_ds = val_ds.cache().shuffle(1000).prefetch(buffer_size = tf.data.AUTOTUNE)
test_ds = test_ds.cache().shuffle(1000).prefetch(buffer_size = tf.data.AUTOTUNE)

# Now after the above step my dataset are kind of optimized for training performa

```

---

```

resize_and_rescale = tf.keras.Sequential([
    # here we did resizing using this api
    layers.Resizing(IMAGE_SIZE, IMAGE_SIZE),

    # here we have done scaling of the image with this api
    layers.Rescaling(1.0/255)
])

data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal_and_vertical"),
    layers.RandomRotation(0.2)
])

input_shape = (IMAGE_SIZE, IMAGE_SIZE, CHANNELS)
# print(input_shape)
n_classes = 3

model = models.Sequential([
    # resize_and_rescale,
    layers.Input(shape=input_shape),

    # Convolutional Layers
    layers.Conv2D(32, kernel_size = (3,3), activation='relu'),
    # layers.Conv2D(32, kernel_size = (3,3), activation='relu', input_shape=input
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, kernel_size = (3,3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, kernel_size = (3,3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),

```

```

        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(64, (3, 3), activation='relu'),
        layers.MaxPooling2D((2, 2)),

    # Neural Network
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(n_classes, activation='softmax'),

])
model.build(input_shape = (32,input_shape))

model.summary()

```

⤵ Model: "sequential\_3"

Layer (type)	Output Shape	Param #
conv2d_7 (Conv2D)	(None, 254, 254, 32)	896
max_pooling2d_7 (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_8 (Conv2D)	(None, 125, 125, 64)	18,496
max_pooling2d_8 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_9 (Conv2D)	(None, 60, 60, 64)	36,928
max_pooling2d_9 (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_10 (Conv2D)	(None, 28, 28, 64)	36,928
max_pooling2d_10 (MaxPooling2D)	(None, 14, 14, 64)	0
conv2d_11 (Conv2D)	(None, 12, 12, 64)	36,928
max_pooling2d_11 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_12 (Conv2D)	(None, 4, 4, 64)	36,928
max_pooling2d_12 (MaxPooling2D)	(None, 2, 2, 64)	0
flatten_2 (Flatten)	(None, 256)	0
dense_4 (Dense)	(None, 64)	16,448
dense_5 (Dense)	(None, 3)	195

Total params: 183,747 (717.76 KB)

Trainable params: 183,747 (717.76 KB)

Start coding or [generate](#) with AI.

```

model.compile(
    optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
    metrics=['accuracy']
)

```

```

history = model.fit(
    train_ds,
    batch_size=BATCH_SIZE,
    validation_data=val_ds,
    verbose=1,
    epochs=25,
)

```

⤵ Epoch 1/25  
54/54 3s 51ms/step - accuracy: 0.7755 - loss: 0.5215 - v  
Epoch 2/25  
54/54 3s 50ms/step - accuracy: 0.8804 - loss: 0.2846 - v  
Epoch 3/25  
54/54 3s 50ms/step - accuracy: 0.9243 - loss: 0.1778 - v

```
Epoch 4/25
54/54 3s 50ms/step - accuracy: 0.9431 - loss: 0.1658 - v
Epoch 5/25
54/54 3s 49ms/step - accuracy: 0.9282 - loss: 0.1757 - v
Epoch 6/25
54/54 3s 58ms/step - accuracy: 0.9525 - loss: 0.1192 - v
Epoch 7/25
54/54 5s 50ms/step - accuracy: 0.9326 - loss: 0.1541 - v
Epoch 8/25
54/54 3s 50ms/step - accuracy: 0.9563 - loss: 0.1162 - v
Epoch 9/25
54/54 3s 50ms/step - accuracy: 0.9512 - loss: 0.1058 - v
Epoch 10/25
54/54 3s 51ms/step - accuracy: 0.9811 - loss: 0.0526 - v
Epoch 11/25
54/54 5s 50ms/step - accuracy: 0.9613 - loss: 0.1101 - v
Epoch 12/25
54/54 5s 51ms/step - accuracy: 0.9801 - loss: 0.0537 - v
Epoch 13/25
54/54 3s 50ms/step - accuracy: 0.9884 - loss: 0.0378 - v
Epoch 14/25
54/54 5s 51ms/step - accuracy: 0.9914 - loss: 0.0187 - v
Epoch 15/25
54/54 3s 50ms/step - accuracy: 0.9768 - loss: 0.0582 - v
Epoch 16/25
54/54 3s 51ms/step - accuracy: 0.9849 - loss: 0.0409 - v
Epoch 17/25
54/54 5s 52ms/step - accuracy: 0.9777 - loss: 0.0560 - v
Epoch 18/25
54/54 3s 51ms/step - accuracy: 0.9757 - loss: 0.0598 - v
Epoch 19/25
54/54 3s 50ms/step - accuracy: 0.9690 - loss: 0.0870 - v
Epoch 20/25
54/54 3s 54ms/step - accuracy: 0.9965 - loss: 0.0143 - v
Epoch 21/25
54/54 3s 56ms/step - accuracy: 0.9922 - loss: 0.0174 - v
Epoch 22/25
54/54 3s 51ms/step - accuracy: 0.9998 - loss: 0.0038 - v
Epoch 23/25
54/54 3s 50ms/step - accuracy: 0.9980 - loss: 0.0130 - v
Epoch 24/25
54/54 3s 49ms/step - accuracy: 0.9974 - loss: 0.0080 - v
Epoch 25/25
54/54 3s 54ms/step - accuracy: 0.9977 - loss: 0.0084 - v
```

◀ ▶

```
scores = model.evaluate(test_ds)
scores
```

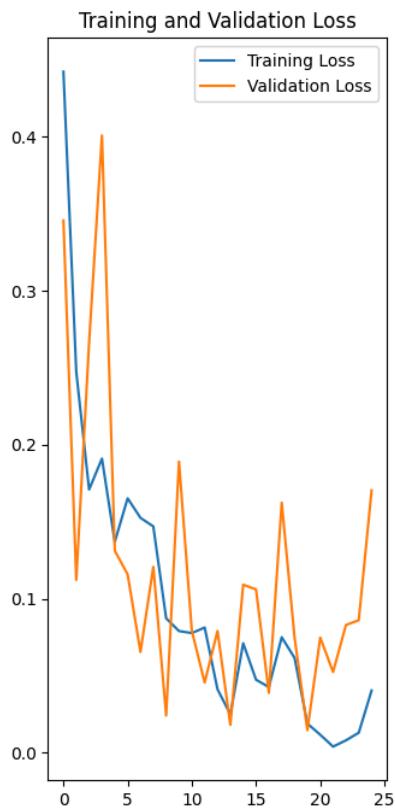
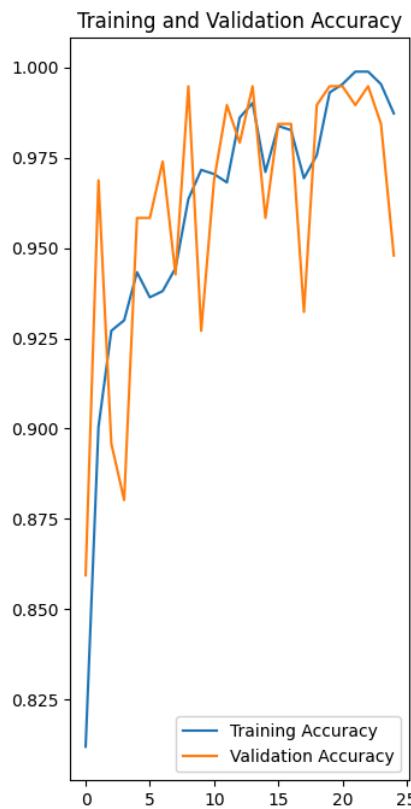
→ 8/8 6s 21ms/step - accuracy: 0.9375 - loss: 0.2064  
[0.19714942574501038, 0.9453125] ▶

```
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
```

```
loss = history.history['loss']
val_loss = history.history['val_loss']
```

```
EPOCHS = 25
plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(range(EPOCHS), acc, label='Training Accuracy')
plt.plot(range(EPOCHS), val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')
```

```
plt.subplot(1, 2, 2)
plt.plot(range(EPOCHS), loss, label='Training Loss')
plt.plot(range(EPOCHS), val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



```
def predict(model, img):
    img_array = tf.keras.preprocessing.image.img_to_array(images[i].numpy())
    img_array = tf.expand_dims(img_array, 0)

    predictions = model.predict(img_array)

    predicted_class = class_names[np.argmax(predictions[0])]
    confidence = round(100 * (np.max(predictions[0])), 2)
    return predicted_class, confidence

import numpy as np
plt.figure(figsize=(15, 15))
for images, labels in test_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))

    predicted_class, confidence = predict(model, images[i].numpy())
    actual_class = class_names[labels[i]]

    plt.title(f"Actual: {actual_class},\n Predicted: {predicted_class}.\n Con
    plt.axis("off")
```

```
1/1 ━━━━━━ 0s 41ms/step  
1/1 ━━━━━━ 0s 44ms/step  
1/1 ━━━━━━ 0s 43ms/step  
1/1 ━━━━━━ 0s 44ms/step  
1/1 ━━━━━━ 0s 41ms/step  
1/1 ━━━━━━ 0s 40ms/step  
1/1 ━━━━━━ 0s 47ms/step  
1/1 ━━━━━━ 0s 41ms/step  
1/1 ━━━━━━ 0s 41ms/step
```

Actual: Potato\_Early\_blight,  
Predicted: Potato\_Early\_blight.  
Confidence: 100.0%



Actual: Potato\_Early\_blight,  
Predicted: Potato\_Early\_blight.  
Confidence: 100.0%



Actual: Potato\_Late\_blight,  
Predicted: Potato\_Late\_blight.  
Confidence: 98.95999908447266%



Actual: Potato\_healthy,  
Predicted: Potato\_healthy.  
Confidence: 100.0%



Actual: Potato\_Late\_blight,  
Predicted: Potato\_Late\_blight.  
Confidence: 99.98999786376953%



Actual: Potato\_Late\_blight,  
Predicted: Potato\_Late\_blight.  
Confidence: 93.75%



Actual: Potato\_Early\_blight,  
Predicted: Potato\_Early\_blight.  
Confidence: 100.0%



Actual: Potato\_Early\_blight,  
Predicted: Potato\_Early\_blight.  
Confidence: 100.0%



Actual: Potato\_Late\_blight,  
Predicted: Potato\_Late\_blight.  
Confidence: 99.81999969482422%



```
model.save("./model.h5")
```

```
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
```